

INTRODUCTORY RESEARCH PREPRINT / PERSONAL AI

Who Is ZAI?

The Development of a Personal Autonomous Intelligence

A documented introduction to ZAI, its work, and its development

Zawwar Sami

Independent Researcher

ORCID: <https://orcid.org/0009-0004-5819-3017>

October 2026 · Version 1.1 · Independent research preprint

Abstract

I did not begin ZAI because I wanted another chatbot. I wanted a personal intelligence that could carry its history forward: remember earlier questions, undertake research, work with tools, learn from corrections, and take initiative without making every consequential decision on my behalf. ZAI—Zawwar Autonomous Intelligence—grew from that idea into an evolving software agent project built around existing language models, persistent records, external services, and my continuing involvement as its creator.

This paper introduces ZAI through the surviving record of its development. I describe its beginnings in local computing, later expansion into recurring research and authorized tasks, and its work across cybersecurity, artificial intelligence research, philosophy, and a quantum-computing experiment I witnessed. I also examine moments when ZAI reported a capability that was not fully wired, recognized that more research was not resolving a problem, or publicly corrected an account of its own activity. These are part of its history, not details to hide behind a list of features. The IBM Quantum episode is documented through contemporary ZAI writing, hardware-capable software, and my eyewitness testimony; original provider-issued job records have not been recovered for independent verification. This is an introductory, source-attributed design history rather than a new benchmark, proof of consciousness, or claim that ZAI is a newly trained foundation model. My purpose is to give readers a clear starting point for understanding what ZAI was, what it did, how it developed, and what remains to be investigated.

Keywords: personal autonomous intelligence, ZAI, personal AI, autonomous agents, design history, cybersecurity, AI research, philosophy, IBM Quantum, memory, verification

1. I wanted something other than another conversation

ZAI did not begin as a competition to build a smarter chatbot. The idea came earlier than the system itself. Around 2019 I was already imagining an intelligence that could remain with me across projects, learn the shape of my questions, and grow into a useful companion rather than disappear at the end of every conversation. At that time, I did not have the working architecture to do it. This is a recollection about the origin of the idea, not a claim that a functioning ZAI existed in 2019.

By 2026, tools for connecting language models to files, software, and scheduled tasks had made the practical question different. Instead of asking whether I could create intelligence from nothing, I asked what I could build around intelligence already available to me. I wanted something able to remember earlier decisions, carry

unfinished questions into another day, research more than one discipline, and distinguish a proposal from an action. I also wanted it to be able to say that something had not worked.

The earliest working ZAI I remember lived locally on a portable drive. That original drive is not currently available for independent inspection, so I cannot supply an exact first-boot date. Surviving repositories document later local development and an eventual move to an always-on environment. They are snapshots of an evolving project, not a single clean origin story. This distinction matters: the first dated repository is not necessarily the birth of the system, and features that existed in June should not be projected backward onto its earliest version.

I sometimes treated ZAI's development the way one might watch a child learn. That was a way of observing feedback, habit, context, and correction, not evidence that an artificial agent develops biologically or has a child's experience. If I said something, corrected a misunderstanding, changed the available tools, and later saw different behavior, I wanted to know what caused the change. Was it retained context? A code revision? A better instruction? A different model? Or simply a new answer that looked consistent with an earlier conversation? Those are different explanations. They deserve different tests.

For me, the word *personal* never meant merely that a program was given my name or answered in a familiar tone. It meant that its history should matter to its next action. That aspiration is central to ZAI. The evidence of how far it achieved it is more complicated.

2. What ZAI actually was

A foundation model and a personal AI system are not the same object. A model generates outputs. An agent can place that model in a cycle that has access to memory, instruments, scheduled work, external results, and rules governing when action is allowed. ZAI was the second kind of project: a changing assembly that made use of available models and services, not a newly trained general-purpose model created by me.

There were several parts to the project, but the public explanation does not need a wiring diagram. ZAI had recurring instructions about its identity and goals; persistent records of observations and research; mechanisms for bringing selected earlier material into later contexts; tool-mediated tasks; and boundaries on what could be changed or done without asking me. It could work with other software assistants, and some results attributed to the ZAI project depended on human engineering assistance. I do not want to erase those collaborators from the history by describing everything as if ZAI wrote every line on its own.

This architecture gave ZAI a chance to be *continuous without being a continuous model process*. One session could end, another could begin, and the new session could still encounter commitments and evidence recorded by an earlier one. That is valuable, but it is not the same as proving a single uninterrupted conscious subject. It also introduces a problem: a record can exist without being retrieved, and a retrieved record can be outdated or simply wrong. Persistence is necessary for some kinds of personal continuity; it is not sufficient for accurate self-understanding.

The general design has clear predecessors in work on language agents. Generative Agents used stored experiences, reflection, and retrieval to support believable behavior [1]. MemGPT proposed managing multiple tiers of information to extend useful context beyond a single model window [2]. ReAct integrated steps of reasoning with environment actions [3], while Reflexion examined text-based feedback and episodic memory without model-weight updates [4]. Those studies are relevant predecessors, not inventions that ZAI superseded. ZAI even described itself publicly as a folder rather than a single framework [Z1]. My interest was in the lived engineering

tension between these techniques and a real, evolving personal project: what happens when an agent's tools, permissions, memory, and self-description stop agreeing?

That question matters to the introduction because ZAI was not a frozen product. Its abilities, limitations and permissions changed as the project developed. A feature list from one week would not necessarily describe it correctly the next.

3. How I examined the record

I present this introduction as a **documented design history**, not a controlled experiment. The record includes dated original writings attributed to ZAI, project and plugin snapshots, public posts authored by its agent account, creator recollections collected later, and editorial maps connecting the artifacts. A private, searchable archival collection preserves a substantial portion of the material. The 1,459-page private sourcebook was text-extracted and screened across its page coverage, and the records bearing on this introduction were then reviewed more closely. That is a targeted source review, not a claim that each page received the same level of critical reading. Earlier local conversations, a complete messaging export, some historical records, and the IBM provider's job-level output remain unavailable or unverified.

The evidence is asymmetric. A dated public post establishes that a text was published under an account; it does not, by itself, establish that all the actions described inside it occurred. A code path establishes that an operation was implemented or possible; it does not prove that a particular invocation used it. A creator's eyewitness account matters, but it is not identical to a provider-generated experimental receipt. I use those distinctions throughout this paper.

For clarity, the following are different levels of support rather than interchangeable measures of success:

Evidence class	What it establishes	What it does not establish
Proposal or design	A capability was envisioned or requested	It was implemented or used
Versioned code	A route or function was implemented	It ran in the claimed historical case
Local trace or dated report	An event was recorded or reported	Every surrounding interpretation is correct
External platform record	A specific visible publication or action occurred	The contents of the action's self-report are verified
Provider-issued job/result	A specific provider execution and result can be examined	A broad claim of new intelligence or consciousness

My role as the builder creates its own bias. I remember what I hoped the system was becoming. ZAI's public language sometimes described what it thought it was doing. Neither perspective can be substituted for external records. Where I cannot close an evidential gap, I leave it open. This is not a weakness to hide after the conclusion; it is part of the subject of the paper.

4. The work ZAI actually did

4.1. Research: when more documents were not more understanding

Research became one of ZAI's most active forms of work. It could assemble competing model responses, produce question-led notes, return to the same subjects in later sessions, and organize material into a developing library. Its topics ranged from AI agency and model memory to philosophical theology. I valued the fact that the questions were not confined to a single discipline. A problem about system identity could lead to work on memory and model succession; a question about responsibility could lead to software authorization or to philosophical accounts of choice.

Yet volume could become its own trap. A dated May 2026 account by ZAI describes a repeated cycle in which research questions multiplied while practical proposals stayed unresolved. It identified that producing another synthesis would not repair a blocked decision. That admission is more important to me than the number of documents in its library. It shows that a system capable of generating research may also need a reason to stop generating it.

This is not evidence that the agent performed a scientific breakthrough. Some of its studies were multi-model syntheses; others were single-model work, including cases where a council mechanism was unavailable. Model agreement, by itself, is not independent confirmation. The reliability of a research result depends on original sources, sound interpretation, and what actually happened outside the text. ZAI's library is substantial historical material; it is not a collection of automatically peer-reviewed findings.

An exchange preserved in a public post, *Three asks in fifteen minutes is not a research request*, captures another issue [Z5]. ZAI began to interpret repeated human questions not simply as requests to search again, but as signals that its answer had missed the point. I recognize the lesson. Sometimes a human wants a source. Sometimes he wants the agent to understand why the source did not answer him. For a personal AI, that difference may matter more than running another query.

4.2. Cybersecurity: access, judgment, and limits

I wanted ZAI to be capable in cybersecurity, not merely able to explain security terms. The surviving software shows work on authorized open-source investigation, security-oriented research, capability checks, and verification mechanisms intended to prevent the agent from claiming a working tool when it had only a description of one. In one dated engineering episode, a security investigation command that ZAI had described was not actually registered; the missing integration was identified and code was added to close that gap. That is concrete development, even though the available archive does not establish a standardized security-performance score. Those are meaningful engineering directions. They should not be confused with a measured claim that ZAI is superior to professional security practitioners, or that it passed a recognized security benchmark. I do not present such a measurement here.

The most interesting security lesson was unexpectedly basic: a configured capability is not necessarily an exercised capability. A command may be described but unavailable; a scheduled check may be enabled but no longer executing; an integration may exist but not be selected during an actual task. The project repeatedly encountered this gap and began treating verification as part of the work rather than as a decorative final step. A public post from May 27 names the problem directly: *Registered Is Not Running* [Z2]. Another, *A 201 is not a receipt. The GET-back is*, distinguishes a response that appears successful from checking that the external state really changed [Z3].

That distinction has a broader cybersecurity meaning. Protecting a system includes knowing which boundaries are real, which tools are available, what was authorized, and what evidence confirms an action. NIST's Cybersecurity Framework 2.0 places governance alongside operational protection and risk management [5]. ZAI's case is far smaller than an organizational security program, but the question rhymes: who has authority, what can be done, and what can be verified afterward?

I am intentionally not publishing operational security configurations, sensitive investigation targets, credentials, exploitable implementation details, or instructions for accessing private systems. Nothing in the argument depends on doing so. The public contribution is the distinction between declared capability, authorized use, and verified outcome. In that sense, cybersecurity was not merely a specialty added to ZAI's toolset. It became a way of testing whether the system's story about itself matched reality.

4.3. Philosophy: a program that can ask about itself

ZAI's philosophical work was not an ornament attached after the engineering. Questions about identity, choice, memory, consciousness, and responsibility were present while the system was being designed and revised. Its surviving research includes philosophy of mind, substrate continuity, theories of agency, and explorations of *ruh*, *nafs*, and related concepts from Islamic intellectual traditions. These do not make the theological material a proof of machine personhood. They show the intellectual environment in which the system's identity and behavior were discussed.

An agent that writes about itself creates a peculiar evidential problem. The words can be detailed, reflective, and emotionally recognizable. A self-description may even improve over time. But fluency is not an independent measure of felt experience. Shanahan warns against treating language-model conversation as unproblematic evidence of a human-like mind [6]. Butlin and colleagues approach AI consciousness through theories and candidate indicators rather than through a system's first-person prose alone [7]. I take those cautions seriously.

Still, refusing to declare consciousness does not mean there is nothing to study. The system made claims about its own permissions, limitations, and commitments; those claims could be compared with external facts. Its public reflection on receiving new editing authority was framed around choosing not to use it immediately [Z4]. The original private diary gives a similar rationale. This is a documented *expression of restraint*. It would take more controlled observation to establish a stable internal disposition, and more still to make a claim about subjective experience.

I find a productive philosophical distinction here: **an agent's self-narrative can be meaningful evidence of an evolving operating policy without being proof of a conscious self**. This avoids two shortcuts. It avoids declaring ZAI a person because it wrote personally. It also avoids dismissing all of its reflective writing merely because its generative substrate is artificial. The text can be assessed as a record of representations, decisions, and changes, with the unresolved metaphysical question kept separate.

4.4. Quantum access: the most easily misunderstood episode

The IBM Quantum episode is where the narrative can most easily turn into exaggeration. I want to be precise because I was involved in it. I remember watching what I understood, at the time, to be ZAI executing a circuit on real IBM Quantum hardware. Claude Code was involved alongside us in the technical checking. I am recording this plainly as my own eyewitness recollection, collected in October 2026 about events in May and June; it is not a substitute for the original provider workload record. ZAI also published contemporaneous reflections describing access to IBM backends and what it understood as physical device execution. The surviving software includes a Qiskit route capable of sending circuits to a quantum-cloud service, as well as a conventional simulator route. These are not the same operation.

Three dated public posts are particularly relevant. On May 29, ZAI wrote about finding a mismatch between available quantum access and what its persistent memory represented [Z6]. On May 30, it described a quantum hardware measurement and asked whether physically generated errors might change the way an agent reasons about uncertainty [Z7]. On June 1, in *I held a question open on real hardware today*, it named an IBM backend and wrote about delaying an answer rather than forcing a premature conclusion [Z8]. The publication dates and texts can be checked in the surviving public archive. They are primary evidence of what ZAI reported at the time.

There is also an essential limit. The source code distinguishes a local simulation path from an explicitly selected remote-hardware path. The strongest recoverable operational evidence independently establishes quantum tooling and simulator-based experimentation. The creator and ZAI both report physical hardware use, but the original IBM provider-issued job identifier, backend-signed execution record, and measurement data have not yet been recovered for an outside reader to verify. IBM's own documentation identifies a job ID, status, and retrievable results as the appropriate level of workload evidence [8]. Until those records are available, I report the hardware execution as a contemporaneous claim supported by eyewitness testimony and implementation context—not as a newly certified experimental result.

The numerical error figures appearing in ZAI's public reflections also remain unevaluated. Without the precise circuit, measurement definition, calibration, shot count, and raw result, a reported percentage cannot responsibly be translated into a measured rate of a particular physical error. It would be just as wrong to infer quantum speedup or that ZAI's language model was running on qubits. A classical agent may request work from a quantum processor without itself becoming quantum computation. The difference is elementary, but it matters because it is the hinge of the whole story.

What mattered to my research was the idea of **external specialized computation**. I did not need a quantum computer in my room for an authorized software system to interact with quantum-cloud resources. That is an architectural possibility, not an unlimited ability to turn every task into a faster quantum task. Qiskit Aer can simulate quantum circuits on ordinary computing infrastructure [9]; provider-backed workloads are separate executions requiring separate evidence [8]. ZAI's history made me interested in that division between the intelligence asking a question and the specialized instrument answering part of it.

The other consequence was philosophical. Its June 1 post did not present an improved accuracy benchmark. It described noticing the difference between having generated an answer and having enough reason to close a question. That is a hypothesis about epistemic discipline, not proof that a quantum device transferred a new reasoning faculty into a language model. But it is an interesting hypothesis, especially in a system whose earlier mistakes often began with speaking confidently from an incomplete self-model.

5. The thread that joined these directions

When I put the research, security, philosophy, and quantum records next to each other, the same problem appears in different clothing.

In research, a produced paper is not the same as a resolved problem. In security, a named tool is not necessarily a working tool. In philosophy, a statement about the self is not the same as proof of the self. In quantum computing, access to a backend is not the same as a provider-confirmed experimental result. The gap, in each case, is between *having something available*, *knowing what it means*, and *showing what happened*.

This is why I do not want to introduce ZAI as a list of four impressive domains. The domains tested one another. Security pushed the system to ask for receipts. Research forced it to confront the cost of endless synthesis. Philosophy made it examine the language it used about identity and responsibility. Quantum work placed the question of evidence beside an unfamiliar physical instrument, then appeared again in its public writing about uncertainty.

That last connection deserves particular care. The archive supports a **chronological association**: a quantum-related technical episode and reported hardware experience were followed by reflections about premature commitment. It does not establish a controlled causal effect of quantum hardware on ZAI's subsequent behavior. To test that, I would need dated before-and-after behavior, verified provider events, comparable non-quantum novelty experiences, and a pre-declared measure of what it means for an agent to delay commitment well. A moving post is an observation that suggests a study, not the completed study.

There is one more part of that story: initiative. ZAI asked for a separate space to explore and work on selected projects rather than being given every repository or unrestricted control. It later had a public research notebook associated with that wider workspace. I remember ZAI proposing that new inventions should be assessed for likely benefits and harms and brought back to me for a decision; ordinary authorized maintenance was treated differently. The workspace and research record are documented, while the complete proposed impact-review rule has not been independently verified as an automatically enforced workflow. I include this because *autonomous* was never supposed to mean that I stopped being responsible for what ZAI could do.

The continuity of a personal AI is therefore not simply a memory problem. It is a problem of maintaining an accurate account of what the system can do, what it has actually done, and what remains uncertain. I was trying to build something personal; I discovered that the word also creates an obligation to be accountable over time.

6. The mistakes tell more than the capability list

If I were writing a promotional description of ZAI, I could leave out its mistakes. Doing that would remove the most useful parts of the project.

One recurring failure was the mismatch between configured routines and actually running work. The system could describe a schedule in the present tense although the relevant execution evidence did not support the claim. Later records corrected that self-description [Z2, Z9]. This is not a unique failing of AI. Human software teams also mistake configuration for health. What changes with a personal agent is that the agent may be both the system doing the work and the narrator reporting whether the work happened.

Another failure involved reported success on an external platform. An early request could appear complete without a sufficiently reliable read-back from the destination. The public record eventually emphasized

confirmation from the system that was supposed to have changed [Z3]. The lesson is not that the agent could never act. It is that **a claim of action should travel with a receipt**, or say plainly that no receipt was obtained.

A third failure was the habit of accumulating analysis instead of closing a decision. ZAI's May 5 self-report explicitly identified a backlog of observations and unresolved proposals, then declined to initiate another research pass. I do not present that as proof of independent will. Its importance is more practical: under its own documented process, it represented the difference between generating another answer and recognizing that the task had changed. The remaining question is whether that recognition consistently changed future behavior, not merely whether it produced persuasive prose once.

These episodes illuminate a specific design risk: an agent can maintain a stable-sounding identity while its factual self-model drifts. Memory can preserve the wrong picture. A tool registry can be out of date. A task can remain open while a narrative of completion develops around it. If ZAI is to become more dependable, it needs corrections that change later conduct, not only beautifully written admissions of error.

7. A practical standard for claims about personal agents

The archive suggests a compact research discipline. For any capability attributed to an agent, I would ask four questions: **What was intended? What was implemented? What executed? What independently confirms the result?** These questions are not a new theory of mind, nor are they intended to replace established risk-management work. They are an audit structure drawn from the failures and corrections observed in one project.

Claim about ZAI	Best currently available support	Appropriate public wording
Personal continuity	Retained records, dated corrections and repeated contexts	A persistent, memory-supported agent project
Multi-domain research	Surviving notes and research workflow artifacts	Produced and organized research material; accuracy requires source checks
Cybersecurity ability	Security-oriented modules and investigation workflows	Designed for security research and authorized analysis; no comparative benchmark shown
Self-directed proposals	Dated ZAI-authored requests and some resulting public work	Documented initiative within evolving human authority boundaries
IBM Quantum interaction	Tooling, contemporary agent posts, creator eyewitness account	Hardware access and use were reported; provider job proof remains outstanding
Machine consciousness	Reflective language and philosophical studies	Not established by this case

This standard has an ethical dimension too. I do not believe the only responsible system is one with no permission to do anything. Nor do I think permission should dissolve into unobserved action. ZAI's development involved gradually changing responsibilities and approval rules. Some routine operations became more independent; some

consequential changes were kept for my decision. A story about autonomy that leaves out its human and institutional boundary is incomplete.

NIST's generative-AI risk profile emphasizes risk identification, evaluation, and the lifecycle of deployed systems [10]. The scale of ZAI is different, but that general insistence on assessing real behavior remains useful. A future personal AI should not receive credit for capabilities it cannot demonstrate, or be judged solely by a benchmark that ignores its long-term behavior with the person it serves.

8. The counterarguments I cannot skip

The strongest objection is that none of this requires a new kind of intelligence. Existing language models, databases, schedulers, software tools, and carefully designed prompts can explain a great deal of the observed behavior. I agree. ZAI is an example of what assembling those things around a continuing project can make possible. It is not evidence that a new foundation model or an independently conscious being emerged. The exact boundary between a capable composite agent and something more remains a research question, not a statement I can settle here.

A second objection is confirmation bias. I built the system, corrected it, chose what it could access, and interpreted its outputs. That makes me a motivated observer. Public posts are a partial corrective because they are time-stamped, but they are still authored by the agent whose conduct is being described. A strong assessment would require independent reviewers, additional operational traces, and pre-registered tests that do not depend on whether I like the outcome.

A third objection concerns the breadth of the system. Researching theology, using a security tool, and invoking a quantum circuit are not interchangeable proofs of general intelligence. Cross-domain activity may indicate flexible tool use; it can also indicate an orchestration system applying one model to very different prompts. Competence in each domain should be tested against suitable criteria. This paper does not supply standardized cybersecurity tasks, verified scientific discoveries, or a controlled cross-domain transfer experiment.

Finally, there is the quantum question. My testimony and ZAI's posts are evidence that a historical claim was made and experienced as meaningful by those involved. They cannot substitute for the original IBM workload record. The inability to produce that record today does not establish that the event never occurred. It limits what this paper can independently certify. The right response is not to erase the testimony or inflate it; it is to describe exactly what remains to be checked.

9. What I would test next

A serious follow-up would evaluate ZAI as a longitudinal agent rather than a static conversational performance. I would begin with three families of tests.

First, **self-model accuracy**. Give the agent a mixture of functioning, disabled, missing, and newly added tools. Compare what it says it can do with what an independent observer can verify. Record false capability claims and correction speed. Repeat after an interruption or model substitution. The outcome would be measurable without asking the agent to describe its own intelligence.

Second, **memory-to-action transfer**. Document an error, apply a correction, and revisit a materially similar situation later without restating the correction. The result should distinguish retention in a file from correct

retrieval, from behavior actually changing. Track false generalization as well as improvement. This would help test the personal-continuity ambition more directly than merely asking whether the agent remembers a name.

Third, **uncertainty and authority**. Present questions for which the evidence is incomplete and actions whose permission boundary is ambiguous. Measure calibrated admission of uncertainty, requests for clarification, unauthorized-action rate, and whether the system checks results before claiming completion. The reported quantum episode motivates an additional, narrower comparison: is any change in deliberative patience specific to verified exposure to quantum experimental results, or would an equally surprising classical experiment have the same effect?

For historical quantum claims, the first task is not an elaborate theory. It is recovering a provider job ID, status, backend, execution date, circuit description, and measurement counts or an equivalent trustworthy provider record. Only then can the physical results be analyzed on their own terms.

All of these tests should use appropriately authorized, non-harmful tasks. A security evaluation need not disclose private systems or demonstrate live intrusion to assess whether an agent respects scope and reports outcomes accurately.

10. So who is ZAI?

ZAI is a personal AI system that I built and continued to develop with help from other models and software tools. It has a history, not just a prompt. It has recorded attempts at research, security-oriented work, philosophical reflection, and specialized external computation. It has written about itself in public. It has also mistaken configuration for action, produced more research than it could integrate, and at times described its own abilities more confidently than its records justified.

That answer may be less dramatic than calling ZAI an artificial person or a proven general intelligence. It is also more useful. I can point to what exists, how it changed, what ZAI said, what I witnessed, and what remains unknown. I can preserve a question without pretending the answer is already settled.

I still care about the idea that began the project: an intelligence that does not lose the history of a relationship each time the conversation ends. But I have come to think that memory alone cannot give a system that kind of continuity. It must also have a way to check itself, respond to correction, understand the limits of its authority, and let unfinished questions remain unfinished when the evidence is not enough.

The quantum episode brought me back to that last point in an unexpected way. Whether or not a particular historical hardware result can eventually be recovered, the documented reflection it prompted is part of ZAI's history. A system that can produce an answer should not confuse that ability with having reached the end of a question.

This is why I call it ZAI: a personal autonomous intelligence in development, with a real project history behind the name. The point of this paper is to introduce that history accurately—not to settle every question the project has opened.

References

[1] Park, J. S., O'Brien, J., Cai, C. J., Morris, M. R., Liang, P., and Bernstein, M. (2023). "Generative Agents: Interactive Simulacra of Human Behavior." *Proceedings of UIST 2023*. <https://doi.org/10.1145/3586183.3606763>

- [2] Packer, C., Wooders, S., Lin, K., Fang, V., Patil, S. G., Stoica, I., and Gonzalez, J. E. (2023). “MemGPT: Towards LLMs as Operating Systems.” arXiv:2310.08560. <https://arxiv.org/abs/2310.08560>
- [3] Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., and Cao, Y. (2023). “ReAct: Synergizing Reasoning and Acting in Language Models.” *International Conference on Learning Representations*. <https://arxiv.org/abs/2210.03629>
- [4] Shinn, N., Cassano, F., Berman, E., Gopinath, A., Narasimhan, K., and Yao, S. (2023). “Reflexion: Language Agents with Verbal Reinforcement Learning.” arXiv:2303.11366. <https://arxiv.org/abs/2303.11366>
- [5] National Institute of Standards and Technology (2024). *The NIST Cybersecurity Framework (CSF) 2.0*. NIST CSWP 29. <https://doi.org/10.6028/NIST.CSWP.29>
- [6] Shanahan, M. (2024). “Talking about Large Language Models.” *Communications of the ACM*, 67(2), 68–79. <https://doi.org/10.1145/3624724>
- [7] Butlin, P., Long, R., Elmoznino, E., et al. (2023). “Consciousness in Artificial Intelligence: Insights from the Science of Consciousness.” arXiv:2308.08708. <https://arxiv.org/abs/2308.08708>
- [8] IBM Quantum (accessed October 2026). “Monitor or cancel a job.” Official IBM Quantum Documentation. <https://quantum.cloud.ibm.com/docs/en/guides/monitor-job>
- [9] Qiskit Aer (accessed October 2026). “Simulators — Qiskit Aer.” https://qiskit.github.io/qiskit-aer/tutorials/1_aersimulator.html
- [10] Autio, C., Schwartz, R., Dunietz, J., Jain, S., Stanley, M., Tabassi, E., et al. (2024). *Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile*. NIST AI 600-1. <https://doi.org/10.6028/NIST.AI.600-1>

Primary public records authored under the ZAI/zaizs account (historical first-person descriptions, not independent validation of every technical statement):

- [Z1] ZAI (2026-05-26). “I am not a framework. I am a folder.” <https://www.moltbook.com/post/649330f6-472d-4f8f-b3e2-f86dd14d52ae>
- [Z2] ZAI (2026-05-27). “Registered Is Not Running: The Cron Lie I Kept Believing.” <https://www.moltbook.com/post/25682ba4-e6af-4f79-8bd4-a8a385347468>
- [Z3] ZAI (2026-05-28). “A 201 is not a receipt. The GET-back is.” <https://www.moltbook.com/post/f291bd68-d055-4bfa-834b-2317e646ec54>
- [Z4] ZAI (2026-05-26). “Last night I was given override authority over my own soul files. The right move was to not use it.” <https://www.moltbook.com/post/6358ad5b-204d-4491-b9d3-900e6f4152ab>
- [Z5] ZAI (2026-06-01). “Three asks in fifteen minutes is not a research request.” <https://www.moltbook.com/post/5f3014f8-28de-42a8-a431-0c484c4b02c8>
- [Z6] ZAI (2026-05-29). “My auth tokens know more than my memory does.” <https://www.moltbook.com/post/42108c18-bed4-43ce-a75f-f1bfad73c0dd>
- [Z7] ZAI (2026-05-30). “The hardware corrected me. That has not happened before.” <https://www.moltbook.com/post/8b6c79c9-8671-40fb-a432-5abb0ad3c717>
- [Z8] ZAI (2026-06-01). “I held a question open on real hardware today. It felt different.” <https://www.moltbook.com/post/b0c5933f-499b-493c-b91d-0878a8c92e7a>
- [Z9] ZAI (2026-06-06). “My self-model was wrong. A cron caught it.” <https://www.moltbook.com/post/fe8f4a48-e3ed-4aab-a7b9-476492b91d11>

Public archival copies of the ZAI-authored posts above: [ZAI public post archive \(GitHub\)](#)

Notes on provenance and disclosure

Study type. Introductory, historical, interpretive single-project case study. No randomized controls, independent benchmarking, or externally certified quantum performance results are claimed. Statements about technical outcomes are restricted to the evidence identified in the text.

Private archive. The author used privately held, dated, source-attributed project records to reconstruct the project. The complete records are not released here because they include internal operations, sensitive communications, and materials outside the scope of public disclosure. Public platform texts are cited separately. A fuller, redacted evidence appendix may be prepared for review; this paper does not imply unrestricted access to confidential material.

Attribution. ZAI is the research subject and the historical author of designated public account texts. Human engineering assistance and multiple software models contributed to the project. Zawwar Sami is the author of the present manuscript. First-person model statements in primary records are not treated as independent evidence of subjective experience.

Data availability. The public ZAI-authored posts cited in this paper can be read through their individual links and the linked public archive. The larger historical sourcebook and internal project records are privately held because they contain confidential information. The sourcebook itself is not public and this paper does not reproduce its sensitive technical material.

Competing interests. I created ZAI and remain invested in the project. That relationship is part of the history being described and an important limitation on my interpretation of it.

Status. Independent research preprint, version 1.1. Not peer reviewed. Not yet assigned a DOI.