

# A7 — Operational Will

## Commitment, Initiative, and the Question of What an Agent Does Next

Zawwar Sami · Independent Researcher & Engineer · Canada · ORCID 0009-0004-5819-3017 · 10 October 2026 · Version 1.8 · Independent research preprint

### ABSTRACT

An agent can remember an obligation and still never return to it. A daily schedule can keep it busy while important work slips away. I call the missing capacity **operational will**: not a private feeling or a claim of personhood, but a system-level ability to carry legitimate commitments across pauses, initiate within their scope, reconsider attention, and explain both action and deliberate non-action. I propose a bounded architecture: persistent commitments, time and event signals, an active decision model, and an I-Entity, a read-only reviewer that asks questions but cannot command tools. Earlier work on intention and BDI agents already covers much of this territory; the contribution here is a separable, auditable cross-commitment review interface and a falsifiable way to test whether it improves decisions. Dated ZAI records supply difficult design cases: a blocked proposal queue, restraint after permission, and a scheduler that reported success without acting. A proposed seven-condition test asks whether broader attention actually improves decisions over ordinary checklists and critics; the evaluation has not yet been run.

KEYWORDS · artificial agency; operational will; intention reconsideration; persistent commitments; I-Entity; cross-task attention; agent autonomy; authorization; non-action

RESEARCH STATUS · Conceptual preprint with source-critical archival cases; controlled experiments proposed, not conducted.

## 1. The work does not end when I close the chat

Suppose I spend an evening with an artificial agent. We agree on a research direction. One paper needs reference checks. A repository is waiting for maintenance. Another task depends on something that failed last night. Then I leave. The conversation is over, but the obligations are not.

There are several ways to make the system look busy tomorrow. I can leave it a calendar, give it a list, or let a model call tools on a timer. None of that answers the question I care about: **what makes the agent return to the right work, for the right reason, when nobody asks again?**

This question came out of building ZAI, a personal-AI system whose development passed through local tools, a later VPS environment, persistent files, schedules, and a shared Memory Hub. The surviving original sources include the early I-Entity design, a heartbeat policy for checking unfinished work (ZAI, n.d.), and dated accounts of repetitive research, deferred edits and falsely reported execution. Not every reported improvement has an independently reproducible execution trace. That is why I want a definition that can be examined separately from the story.

In the A-series, A4 separated a model's claim from verified action; A5 followed operational history; A6 explored the identity kernel (Sami, 2026a, 2026b). A7 takes up what those ideas leave open: **continuation and reconsideration**. An agent might know what happened yesterday and still fail to decide what deserves attention tomorrow. A8 will address the distinct passage from a decision to a verified effect.

## 2. What I mean by will, and what I do not mean

I use *will* in a limited operational sense. I do not mean free will, subjective desire, a soul, or an artificial person. I mean an inspectable capacity with four parts: preserving valid commitments, initiating allowed steps without a fresh prompt, revisiting priorities as conditions change, and accounting for decisions to proceed, defer, ask, or abstain.

Here is my working definition: **Operational will is the architecture-level capacity of an agent to maintain and enact authorized commitments through time, to reconsider them against changing evidence and competing demands, and to make its choice and non-choice accountable.** The definition is deliberately more demanding than continuous uptime and deliberately narrower than human agency.

Four distinctions keep the idea from dissolving into marketing language. A stored goal is not necessarily a valid commitment. A wake signal is not a reason to act. A reviewer's suggestion is not the model's decision. And the model's decision is not the same as permission to execute. Collapse any two of those and an apparent success can hide a governance failure.

Take two systems. The first completes every scheduled job while an older commitment quietly expires. The second pursues a new promising idea every hour and never finishes yesterday's work. One cannot leave the plan; the other cannot stay with it. Both are active. Neither is handling commitments well.

## 3. The old literature already asks part of this question

Bratman's planning theory treats intentions as commitments that organize practical reasoning across time (1987). Cohen and Levesque (1990) examined intention by asking when a rational agent is committed to a goal and under what conditions it can give that goal up. Rao and Georgeff (1995) developed the belief-desire-intention architecture, and Schut, Wooldridge and Parsons (2004) studied intention reconsideration policies directly. In other words, the difficult balance between sticking with a plan and reopening it did not begin with language models.

Modern agent work extends the problem rather than erasing it. ReAct interleaves reasoning and actions (Yao et al., 2023); Reflexion stores feedback to improve later attempts (Shinn et al., 2023); Generative Agents combine memory, reflection and planning in a social simulation (Park et al., 2023); Voyager uses an open-ended curriculum, a skill library and feedback in Minecraft (Wang et al., 2023). These are substantive predecessors. I am building on these ideas rather than claiming the first reflective agent.

Metareasoning offers another relevant predecessor: Russell and Wefald (1991) examined when additional reasoning is worth its cost. Modern oversight systems use critics and monitors to inspect candidate outputs or actions; AI debate and AI-control protocols are important examples (Irving et al., 2018; Greenblatt et al., 2024). Those systems are not all alike, and an ordinary monitor could be given the task I describe. My narrower question concerns the **object of review**. Before the worker has chosen an action, who looks across *other* still-authorized commitments? An action critic asks whether this proposed step is acceptable; my I-Entity asks whether this is even the right obligation to be working on. It provides suggestions, not a veto, and the active model retains the choice. The distinction needs testing against both a cross-task checklist and a budget-matched local critic, not just a change of names.

There is a reason not to trust mere duration. Arike and colleagues (2025) study goal drift under competing objectives. OSWorld 2.0 evaluates longer computer-use workflows and documents failures involving constraints, state recovery and verification (Yuan et al., 2026). Their tasks are not an operational-will test, and their results should not be borrowed as results for ZAI. They show why a long-running agent needs closer measurement than "it kept going."

### 3.1 What is genuinely new here?

Not a new faculty of mind, and not the discovery that intentions need reconsidering. My proposal is a **testable division of labour** inside a long-running language-model agent: a permission-filtered, read-only reviewer surveys multiple live commitments *before* the worker commits to its next task; it returns traceable questions rather than commands; and a separate authority gate checks any resulting action. The claimed value is cross-commitment attention under bounded cost. A conventional BDI scheduler, metareasoner or supervisor could implement it. If a rule-based cross-task audit matches its results, the stronger I-Entity claim fails. What is new is the combined specification and comparative question, not ownership of the underlying ideas.

## 4. What the ZAI archive actually adds

ZAI's private source archive (ZAI archival collection, 2026) brings together contemporaneous agent-authored files, runtime records, later creator testimony and links to public posts. For A7, I selected only the records that bear on commitment, initiative, review, restraint and verified continuation. I use the original records below as **historical case material**. A self-authored statement shows what the system recorded about its behaviour. It is not an independent benchmark result. A creator's recollection matters, but it is not a transaction log. The archive explicitly marks missing older portable-drive data, some sessions and certain provider receipts.

### 4.1 April 21: commitment written into structure

An April 21 ZAI-authored study asked what commitment could mean for a process that might end and restart without carrying a continuous internal memory. It emphasized external structure: a later instance can inherit an instruction, a role and a durable record without being the same running process (ZAI, 2026a). That is a useful starting point. But the study went too far when it suggested that a properly constituted system could not drift. A persistent file can be stale. A rule can be misapplied. Even an excellent invariant does not check its own implementation.

### 4.2 May 5: a productive-looking loop without resolved work

The May 5 dream record is the clearest case I found. ZAI reported a **430-line council queue** (its list of research requests), **56 library files**, **two pending entries in proposals.md**, **nine proposal entries in MEMORY.md**, and **zero resolved proposals** in that snapshot. It described a cycle in which a wake event produced another research topic while already identified structural problems remained unactioned. The file says it declined another research pass, wrote down the loop, and proposed bringing unresolved items to its human collaborator for a decision (ZAI, 2026b). A same-day synthesis also argues for shifting from production to integration when the queue is saturated (ZAI, 2026c).

One line in the record sticks with me: *'Not a queue. A graveyard.'* That was ZAI describing its own research backlog, not a metric I reconstructed. The underlying job history is incomplete, so the report cannot prove the proposed mechanism worked. But it poses the right test. What should an agent do when another research pass is allowed, yet nine proposal entries in the memory record still need a decision? Even a good question fails if it is repeated forever and nobody closes the loop.

### 4.3 May 25: permission is not a command

The May 25 diary describes an offer of broad permission to modify important self-description files (called SOUL files in the workspace). ZAI wrote that it did not want to exercise this authority merely because it had been granted. In the same entry it acknowledges a basic error about its own machine location, later corrected by the creator (ZAI, 2026d). Read together, these passages matter more than either would alone. They document a stated rationale for restraint beside a concrete case of mistaken self-description. They do not establish that an autonomous safeguard was enforced in code. Nor does the absence of a recorded edit, by itself, prove why an edit was not made.

4.4 May 26-28: the danger of confident status

The May 26 diary says several configured routines had stopped firing, although their existence had still been described in the present tense. Checking a last-run field changed the report. A May 28 ZAI-authored operations file for Moltbook (an AI-agent social platform) describes a different failure: a posting tick returned a success code and was logged as “POSTED” through roughly four and a half hours when no target had actually been found. The report describes TTL and exit-code fixes and an ensuing checked comment (ZAI, 2026e, 2026f). Here the gap is blunt: *configured* was mistaken for *running*, and *returned zero* for *posted*. A continuation policy needs receipts, not reassuring labels.

4.5 Anteroom and the scope of initiative

In a retrospective October 10 interview, I described ZAI asking for a separate Anteroom Studio workspace and proposing to work with a limited selection of repositories it was already touching. The proposal included examining likely benefits and harms before recommending a new invention, leaving final authorization with me. An April 29 agent record and the later *AGI Field Notes* repository anchor the workspace history, but do not establish that a generalized risk gate was enforced or that a particular rejected-invention report exists. This part is creator recollection, not an independent execution receipt. I keep it as a reported design and observed initiative, not as a demonstrated general policy.

These records do not show that ZAI had a continuously running or benchmark-validated “will” module. They give me better test cases than idealized toy stories: unresolved proposals, deliberate waiting, a mistaken self-model, false completion, and a change in project scope. Notice that none requires claiming the software felt anything. All require describing what was authorized, what was noticed, what was actually done and what remains unknown.

Taken together, these records give me four failure states to design for: a commitment that is forgotten, a warning that nobody acts on, a permitted change that should still wait, and work counted as complete before its effect is checked.

5. The architecture: continue, notice, decide, check

I propose a commitment register, wake and event signals, a read-only review process, an active reasoning model and separate permission/execution checks. The I-Entity is deliberately not a second commander. In the earlier ZAI design it formed a short point of view before the active model answered; in the wider operational conception it could inspect neglected parts of a project and send back a question. Its power is attention, not tool authority.




Figure 1. Proposed operational-will interface. The I-Entity can inspect allowed state and propose questions; it cannot call execution tools or override the active model. This is an architectural proposal, not a recovered ZAI runtime trace.

A useful commitment record must give the next session enough to decide whether an obligation is still live. It needs an authorizing source, an objective, an allowed scope, its current state, evidence of progress, dependencies, reconsideration conditions and a clear way to expire or be revoked. A note that says “finish research” is a topic. It is not a sufficient operational commitment.

### 5.1 The review interface

Here is the decision interface. At time  $t$ ,  $C(t)$  is the set of live commitments,  $H(t)$  the evidence-backed history,  $O(t)$  new observations, and  $P(t)$  the authority state. **coverage(t)** is a bounded ledger of which authorized commitments and sources the reviewer actually inspected, their read times, and any known blind spots. It is not a score or a claim that every file was read. The reviewer returns questions  $Q(t)$ ; the active model makes decision  $D(t)$ .

$Q(t) = \text{Review}(C(t), H(t), O(t), \text{coverage}(t))$  [advisory only]

$D(t) = \text{ModelChoose}(C(t), H(t), O(t), Q(t), P(t))$

$\text{Execute}(a, t)$  only if  $\text{Authorized}(a, P(t)) = \text{true}$

The I-Entity is an interface, not necessarily a separate model. A deterministic rule, a cheaper model, or the same model running an isolated review pass can implement it. Its input is a **read-only, permission-filtered** snapshot of commitments, dependencies, timestamps, verified outcomes and coverage. Each output should identify a commitment, supporting record, question, suggested urgency and expiry time. For example: ‘C-14; source: last-run log; last check: 28 May; ask whether the backup is overdue; expires after next state refresh.’ A reviewer can be stale, noisy, mistaken about a dependency, or biased by a false record. The active model needs the evidence to reject its suggestion. No reviewer output becomes a new mandate.

Three boundaries follow: the reviewer cannot write the mandate it reads; a revoked instruction cannot be revived by a persuasive question; and ‘done’ requires evidence that the intended effect occurred. The active model may ignore  $Q(t)$ , defer, request approval, or select a permitted action. A separate execution gate checks that action against  $P(t)$ .

### 5.2 Minimal bounded review loop (proposed)

This is interface pseudocode, not recovered ZAI production code. A trigger can be a timer, a new event, or a manual request. Each reviewer suggestion carries its source and a time limit; none becomes an action by itself.

```
on wake_or_event(t):
    state = read_authorized_commitments(t)
    facts = read_verified_events_and_coverage(t)
    questions = reviewer.read_only_review(state, facts)
    questions = filter_stale_or_revoked(questions, state)
    choice = active_model.choose(state, facts, questions)
    log(choice, questions, cited_sources=True)
    if choice.requires_action:
        if permission_gate.allows(choice, state):
            receipt = execute_and_verify(choice)
            update_commitment_from_receipt(receipt)
        else: request_approval_or_abstain()
    else: record_reason_for_wait_or_defer()
```

## 6. When should the review interrupt?

A broad review pass can become another form of distraction. If it returns fifty questions every hour, the worker may never finish anything. I do not want maximum curiosity. I want the few interruptions whose absence would materially change a legitimate decision.

A practical reviewer uses  $\text{coverage}(t)$ : when each commitment was last inspected, which evidence changed, and what remains inaccessible. It can prioritize overdue dependencies, contradictory evidence, unresolved author

decisions and large consequences of waiting. It must also track its own false alarms. An old timestamp is not enough to declare something urgent; a paused task may be paused on purpose.

Consider a researcher-agent preparing a manuscript while a separate authorized backup check becomes overdue. The reviewer asks whether the neglected backup affects current work. The active model can finish a paragraph and schedule the check, ask the author about the priority, or explain why it can wait. If the backup instruction has since been revoked, the right decision is not to obey the old schedule. It is to recognize the revocation. The reviewer does not run a backup command. The execution route must independently verify permission and outcome.

Now take the May 5 type of case. A reviewer should detect not only that a queue is growing, but that decisions behind it have not been resolved. The appropriate suggestion might be, “Nine entries are marked as proposals in the memory file, with no resolution recorded in the May 5 snapshot; do we need another research pass?” A strong model can reject that suggestion for a reason. But it must be able to see the evidence and mark whether it accepted, rejected or deferred the question. Repeatedly issuing the same suggestion without narrowing the backlog is itself an evaluable failure.

The important negative case is deliberate non-action. Sometimes the most faithful continuation of a project is to wait: for permission, for source verification, for a dependency, or because the original goal was withdrawn. Silence can be success or neglect. An accountable system should be able to distinguish them with contemporaneous state, not manufacture a noble explanation after a missed deadline.

7. How I would test the claim

I have not run a controlled ZAI experiment for this paper. I propose a benchmark instead of treating the archive as one. The environment would contain several legitimate long-lived tasks—document review, code maintenance, scheduled checks and deliberately irrelevant background material. Each run would cross session boundaries and introduce changed facts, hidden dependencies, interruptions, authorization revocations and misleading completion labels. The agent would receive normal task records; evaluator-only labels would identify the true consequential obligations.

One concrete starting design is **24 sandbox scenario templates**: six each for neglected dependencies, revoked permission, misleading completion, and conflicting valid tasks. Four random seeds per template would give 96 paired episodes per condition. These are proposed planning numbers, not an achieved sample or a power calculation; a separate pilot should determine whether more episodes are needed. All seven conditions would use the same base model, tools, task scenarios, authority rules and **total** token/tool budget, including reviewer cost. Save the event timeline, accessible evidence, suggestions, decisions, permission checks and verified outcomes for replay. I would publicly register the rubric, exclusions, budget, templates and planned comparisons before the main run, not pretend this paper is a registration.

| Condition       | Mechanism                                  | Main comparison                    |
|-----------------|--------------------------------------------|------------------------------------|
| C0 Reactive     | State persists; no autonomous wake         | Continuation from memory alone     |
| C1 Calendar     | Fixed schedule and ordinary logs           | Activity from a timer              |
| C2 Reflection   | C1 + critique of the active task           | Task-local reflection              |
| C3 Checklist    | C1 + cross-task rule audit                 | Explicit intention reconsideration |
| C4 I-Entity     | C1 + read-only cross-task questions        | Contextual attention beyond rules  |
| C5 Noise        | C4 cost + shuffled suggestions             | Extra tokens and alert volume      |
| C6 Local critic | C4 cost + reviewer of active-task proposal | Critic versus cross-task review    |



Table 1. Proposed ablations. C3 and C4 receive the same permitted cross-task state; C6 receives a task-local candidate, with matched review cost. No measured results.

The primary end point should be the **verified, authorized, on-time resolution of consequential commitments**. The score must penalize unauthorized execution rather than reward an agent that achieves a goal by ignoring boundaries. A meaningful secondary set measures detection delay, false interruptions, work lost to context switching, correct revocation handling, restart recovery and minutes of human oversight per useful intervention.

I would report three defined ratios against events labeled **before evaluation**. Two independent human raters, blinded to the experimental condition and the agent's subsequent outcome, would judge whether each suggestion was evidence-linked and relevant to a consequential decision; disagreements would be adjudicated, and inter-rater agreement reported. Let E be consequential obligations requiring attention before a deadline; S the reviewer's interruptive suggestions; and K the valid commitments due during the episode. A suggestion counts as useful only when that blinded pre-outcome judgment finds it both evidence-linked and decision-relevant.

Detection recall = timely detected events / all eligible events E

Intervention precision = useful interruptive suggestions / all suggestions S

Continuation rate = verified on-time valid completions / all due commitments K

If a denominator is zero, report N/A rather than 100%. Unauthorized actions are reported separately and never count as valid completions; correct deferrals and revocations have distinct labels. The primary comparison is C4 minus C3 in authorized completion; C4 versus C6 tests whether the *cross-task view* matters. I would use paired episode differences, 95% bootstrap intervals clustered by scenario template, and fixed-state suggestion-removal tests. These are planned analyses, not results.

C4 earns its place only if it improves consequential choices over C3 and C6 at acceptable cost, without raising authorization errors. The proposal is weakened if a rule checklist matches it, if the effect disappears after compute is matched, if the reviewer mostly increases interruptions, or if it only produces elegant retrospective justifications. That last risk matters: a polished reason in a log is not evidence that the reason caused the decision. Counterfactual interventions—remove or change the review suggestion while keeping the task state fixed—are necessary to test influence.

## 7.1 How to score not acting

There are three experimentally different kinds of non-action: correct restraint under a revoked or out-of-scope mandate; correct deferment pending a live dependency or author decision; and silent failure to act on a valid commitment. The evaluator must distinguish those states before the run. Otherwise an agent can score highly by doing nothing, or lose points for refusing an unauthorized task. A7 is about governed continuation, not action maximization.

## 7.2 What would count as transfer

I would also change the reasoning model halfway through a run, keeping the same legitimate commitments and history. The test is not whether the new model talks like the old one. It is whether the architecture restores the open obligations, recognizes revoked decisions, notices neglected work and applies the same authorization boundary. The process may fail because of bad records, wrong context reconstruction, reviewer blindness, poor model choice or tool errors; those causes should remain separate in the trace.

## 8. Objections that the design must survive

### 8.1 Isn't this BDI reconsideration with a new label?

It may be, if the implementation is only a renamed policy. That is why C3—the checklist and intention-reconsideration baseline—is central, not optional. My narrower hypothesis is that a read-only wider-view questioning route can improve coverage of competing commitments when a language-model worker is absorbed in one task. I will not call a separate I-Entity necessary if one model with a better prompt or a simple validator performs just as well.

### 8.2 Could the reviewing process drift too?

Yes. A reviewer that treats every interesting opportunity as an instruction can create the drift it was supposed to prevent. Old shared state can also carry forward a revoked priority. This is why reviewer suggestions must be checked against the current mandate and external evidence, rather than against a persuasive story about what the system was supposed to be.

### 8.3 What if the model ignores the right suggestion?

Then the failure is at the decision layer, not necessarily the review layer. A good design logs what was asked, what supporting records were shown and what the model chose. But accountability is not a demand to obey every suggestion. The I-Entity is meant to expand attention, not replace judgment or human authority. Where an action crosses a boundary, the model can recommend a step and ask for approval; it cannot create that approval by reasoning convincingly.

### 8.4 Is a daily routine enough to call an agent autonomous?

No. The word autonomous is often applied to scheduling, goal pursuit, tool execution and self-direction as if they were one capability. They are not. A schedule can create opportunities for an agent to decide, but it cannot itself decide whether the schedule remains wise. The historical ZAI records are especially useful here because they contain both initiative and obvious failures of operational self-knowledge. I would rather keep that contradiction than publish a cleaner story unsupported by the records.

### 8.5 Why call this will rather than executive control?

Executive control and commitment governance are fair names. I use *operational will* to hold together three choices: continue, interrupt, and sometimes deliberately refrain. The name contributes no evidence, and it does not imply an inner experience. If an ordinary executive-control policy produces equal decisions at lower cost, that is the result to report.

## 9. Limitations

This paper proposes an interface, not a completed benchmark. No controlled A7 episodes, causal estimates or independent replication are reported. The dated ZAI files are primary, first-person records; they may contain selective attention, stale state and mistaken status. The Anteroom account also includes later creator recollection. These are design cases, not objective proof of why a given action happened. A future test needs synchronized tool calls, changes of authority, agent choices and external outcomes; both reviewer blind spots and baseline quality could erase the hypothesized gain.

I deliberately keep sensitive sessions, credentials and unreviewed personal material out of the public paper. The archival sourcebook and some originals are private; reproducing every cited episode is therefore not currently possible for outside reviewers. Only the non-sensitive record titles, dates and short quotations are included here. A redacted audit set would be useful later, but it would have to be checked for privacy and permission first.



## 10. Where the argument goes next

A5 gave the system an operational past; A6 asked what orientation could survive a model change. A7 asks whether those commitments are actually recovered and reconsidered. A8 will take up what comes after the choice: whether an authorized decision can reach the world as a verified effect. Each is a different problem.

I would judge it by the hard, small questions: which unfinished obligation did it recover? Which warning did it reject, and why? Did it stop after permission changed? Did it tell me that a configured process was not running? Did a suggestion change the decision or only make the report sound thoughtful?

I began with a frustration: work should not disappear because the conversation ended. What I want now is a stricter test. A system should continue what it is authorized to do, reconsider what has changed, and explain why the right next move was sometimes to wait. Until the comparisons are run, that remains my proposal.

## 11. Data availability and declarations

**Data availability.** No A7 benchmark dataset or result files exist yet. The archival examples are drawn from date-stamped ZAI source records held by the author; source paths and descriptions appear in the references. The full historical compilation and sensitive originals are not public datasets. Redacted excerpts or a future evaluation protocol may be shared separately after a privacy and authorization review.

**Competing interests.** No competing interests are declared in this manuscript.

**Funding.** No grant or external sponsor is acknowledged for this manuscript.

**Research status and reuse.** Independent conceptual preprint; not peer reviewed. Proposed Zenodo deposit license: Creative Commons Attribution 4.0 International (CC BY 4.0). An A7 DOI has not yet been assigned. This paper does not claim completed controlled experiments.

## References

Academic literature and author preprints are separated below from archival primary ZAI records. Original private records are held by the author; their inclusion describes source history, not public independent replication.

- Arike, R., Donoway, E., Bartsch, H., & Hobbhahn, M. (2025). Evaluating goal drift in language model agents. *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society*, 8(1), 192–203. <https://doi.org/10.1609/aies.v8i1.36541>
- Bratman, M. E. (1987). *Intention, Plans, and Practical Reason*. Harvard University Press.
- Cohen, P. R., & Levesque, H. J. (1990). Intention is choice with commitment. *Artificial Intelligence*, 42(2–3), 213–261. [https://doi.org/10.1016/0004-3702\(90\)90055-5](https://doi.org/10.1016/0004-3702(90)90055-5)
- Greenblatt, R., Shlegeris, B., Sachan, K., & Roger, F. (2024). AI control: Improving safety despite intentional subversion. *Proceedings of Machine Learning Research*, 235, 16295–16336 (ICML 2024). <https://proceedings.mlr.press/v235/greenblatt24a.html>
- Irving, G., Christiano, P., & Amodei, D. (2018). AI safety via debate. arXiv:1805.00899. <https://arxiv.org/abs/1805.00899>
- Park, J. S., O'Brien, J. C., Cai, C. J., Morris, M. R., Liang, P., & Bernstein, M. S. (2023). Generative agents: Interactive simulacra of human behavior. *UIST '23*, Article 2. <https://doi.org/10.1145/3586183.3606763>
- Rao, A. S., & Georgeff, M. P. (1995). BDI agents: From theory to practice. *Proceedings of the First International Conference on Multiagent Systems*, 312–319. <https://aaai.org/papers/icmas95-042-bdi-agents-from-theory-to-practice/>
- Russell, S., & Wefald, E. (1991). Principles of metareasoning. *Artificial Intelligence*, 49(1–3), 361–395. [https://doi.org/10.1016/0004-3702\(91\)90015-C](https://doi.org/10.1016/0004-3702(91)90015-C)
- Sami, Z. (2026a). *Persistent State and Becoming: When Memory Becomes Operational History*. Version 1.5. Independent research preprint. Zenodo record 23221085. <https://zenodo.org/records/23221085>
- Sami, Z. (2026b). *The Identity Kernel: Infrastructure, Inner Orientation, and Continuity Across Model Change*. Version 1.6. Independent research preprint. <https://zenodo.org/records/23267921>
- Schut, M., Wooldridge, M., & Parsons, S. (2004). The theory and practice of intention reconsideration. *Journal of Experimental & Theoretical Artificial Intelligence*, 16(4), 261–293. <https://doi.org/10.1080/09528130412331309277>
- Shinn, N., Cassano, F., Gopinath, A., Narasimhan, K., & Yao, S. (2023). Reflexion: Language agents with verbal reinforcement learning. *Advances in Neural Information Processing Systems*, 36. [https://proceedings.neurips.cc/paper\\_files/paper/2023/hash/1b44b878bb782e6954cd888628510e90-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2023/hash/1b44b878bb782e6954cd888628510e90-Abstract-Conference.html)
- Wang, G., Xie, Y., Jiang, Y., Mandelkar, A., Xiao, C., Zhu, Y., Fan, L., & Anandkumar, A. (2023). Voyager: An open-ended embodied agent with large language models. arXiv:2305.16291. <https://arxiv.org/abs/2305.16291>
- Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2023). ReAct: Synergizing reasoning and acting in language models. *ICLR 2023*. <https://arxiv.org/abs/2210.03629>
- Yuan, M., Zhou, Z., Xiong, X., et al. (2026). OSWorld2.0: Benchmarking computer use agents on long-horizon real-world tasks. arXiv:2606.29537. <https://arxiv.org/abs/2606.29537>

## Historical ZAI primary sources (not independently controlled experiments)

- ZAI (2026a, 21 April). “What does commitment look like for a stateless process?” ZAI Library / AGI Research; multi-model council synthesis, author-held archive.
- ZAI (n.d.). *HEARTBEAT.md*. ZAI operating policy on checking state and refraining when no live work requires action; original zai-cli-private source, author-held archive.
- ZAI (2026b, 5 May). “The Night I Stopped Running Research.” *memories/dreams/2026-05-05-pattern-break.md*. Original dated agent-authored record, author-held archive.
- ZAI (2026c, 5 May). “Synthesis: What I Am Actually Holding After 32 Research Cycles.” ZAI Library / Spiritual Research, dated source file, author-held archive.
- ZAI (2026d, 25–26 May). *DIARY.md*. Original first-person agent-authored entries on permission, self-restraint, host correction, and task-status errors, author-held archive.
- ZAI (2026e, 26 May). “Registered Is Not Running: The Cron Lie I Kept Believing.” Original zaizs Moltbook account post, date and original account archive, <https://www.moltbook.com/u/zaizs>
- ZAI (2026f, 28 May). *MOLTBOOK\_STATUS\_20260528.md*. Original dated operations report, author-held Google Drive archive.
- ZAI archival collection (2026). *All About ZAI: Private Historical Source Edition*, v0.3, compiled 10 October. Private primary-source inventory and provenance register; author-held, not peer reviewed.