

When Is an AI Personal?

Memory, Continuity, and the Authority to Revise

Zawwar Sami

Independent researcher, Canada

Abstract

An AI assistant can remember a person in considerable detail while leaving that person with little control over how the record is used. This paper distinguishes personalization from a stronger, normative sense of personal assistance. Its central claim is that when an assistant relies on accumulated personal history to interpret instructions or act on someone's behalf, the person must retain effective authority to contest and revise that history's practical force. The argument begins with a shared memory system developed for use across AI clients, then examines a case in which an assistant mistakes an old project decision for a standing preference. Work on extended cognition helps explain why such records can matter to agency; work on generative cognition complicates the analogy with a passive notebook. Correction, traceability, and usable transfer are defended as connected supports for continuing authority, rather than as a checklist sufficient for autonomy. The resulting account explains why possessing a database, preserving an assistant's personality, and accumulating more memory can each fall short of making an AI meaningfully personal.

Keywords: personal AI; memory; human agency; extended cognition; autonomy; AI assistants

1. The problem with starting again

The practical problem that led me to build a shared memory hub was ordinary. An assistant in one place did not know what had already been settled in another. A decision made while working on a server had to be explained again in a browser. Moving between tools meant carrying the background along by hand. Each assistant might be capable, but the work between them remained fragmented.

The project that came out of this, zai-memory-hub, places memories and decisions in a shared database that different AI clients can consult. Its public documentation describes records of decisions, their reasons, and the decisions they supersede. It also provides for inspecting records and transferring them to another installation (Sami, n.d.). The project supplies the starting problem for this paper. It is not offered as an experiment demonstrating improved reasoning or as a complete implementation of the account defended below.

In an earlier essay, “Two AIs,” I described a distinction between transactional use of AI and an ongoing relationship supported by personal memory (Sami, 2026). That distinction needs greater precision. A hosted assistant can retain history; a system on a privately controlled server can misinterpret it. Location alone does not settle whether the relationship is personal in the relevant sense. Nor does continuity automatically benefit the person whose life is being recorded.

Consider the difference between an assistant that remembers an old decision and one that keeps treating it as binding after the reasons for it have disappeared. Both have memory. Only one is

responding appropriately to change. A record that once saved the user from repeating an explanation can later make the same user struggle to be heard.

The question, then, is not simply how much an AI should remember. It is what authority remembered material should have. The claim defended here is limited: where an assistant uses a person's history as a basis for acting on that person's behalf, its use of that history should remain effectively answerable to the person. A user should be able to discover, contest, and revise the practical commitments attributed to them. Continuity that defeats those abilities can preserve a profile while undermining the agency it was supposed to support.

2. Personalization and personal authority

“Personal” can mean several things. A response may be personalized because it mentions the user's occupation or adopts a preferred tone. A service may be personal because only one person uses the account. Neither fact tells us who can determine what the assistant takes the person's past to authorize.

This paper uses personal assistance in a stronger, normative sense: assistance organized around a person's continuing authority over the purposes for which their history is used. This is a proposed standard for a relationship of assistance, not a claim about every ordinary use of the word. A recommendation system may accurately predict what someone will click without satisfying that standard. Prediction and representation have different demands. Knowing what someone is likely to do does not, by itself, entitle a system to decide what should be done on their behalf.

Gabriel et al. (2024) characterize advanced assistants in terms of planning and executing actions for users, and examine questions of alignment, influence, trust, and privacy. That framing makes the distinction between action for a user and mere behavioral prediction especially relevant. The present argument isolates one issue within that wider field: the authority an assistant assigns to accumulated personal context.

Suppose a researcher rejects an expensive software service because a small project cannot justify the cost. An assistant stores the decision as “prefers free tools.” A second assistant later uses this summary to exclude a paid service from a different, funded project. The recommendation may look attentive. It may even come with a friendly reminder that it respects the researcher's preferences. Yet a conditional decision has become a general rule.

There are two errors here. The first is representational: the stored summary loses the reason and scope of the original decision. The second concerns authority: the assistant allows its interpretation of an earlier choice to govern a later one. Correcting the first sentence in a database will not resolve the second error if derived summaries or other connected clients continue to treat the old inference as binding.

The user need not inspect every retrieval before receiving advice. That would defeat much of the purpose of assistance. But when a recommendation depends materially on a personal attribution, the user should have a workable way to ask why that attribution is being made and to change its standing. A statement such as “that was only for the previous project” must be capable of altering subsequent assistance. A system that repeatedly acknowledges the correction while preserving the old practical restriction has not given the user effective control.

This is why apparent familiarity is a poor test. An assistant can sound as though it knows someone while making it difficult for that person to depart from its account of them.

3. Why control of a record matters to agency

The argument can be stated through three premises. First, acting on someone's behalf gives their present, informed directions presumptive authority over an inferred account of what they used to want. Second, personal memory can influence which options an assistant presents and which actions it proposes or performs. Third, an authority that cannot be exercised against a materially influential attribution is ineffective in that part of the relationship. Together, these premises support a requirement of effective correction whenever remembered context acquires practical authority.

The first premise does not imply that every current instruction must be obeyed. Assistants may have independent reasons to refuse an action, and users can be mistaken about the world. The narrower point is that an assistant cannot justify overruling a person's clarification merely by appealing to its own profile of that same person. If a refusal is justified by another obligation, it should be defended on that basis. “Your history says otherwise” is not an adequate answer to a competent user's explanation that a preference has changed.

Clark and Chalmers (1998) offer a useful way to understand the significance of external records. Their case of Otto's notebook concerns information that is reliably available and routinely used to guide action. They argue that, under the relevant conditions, a resource outside the body can contribute to cognition itself. Their discussion also attends to trust and endorsement; not every accessible inscription becomes a person's belief.

The present argument does not require accepting their stronger claim about where the mind ends. Even if a database remains wholly external to the mind, its role in selecting options can make its governance important. Imagine that the researcher in the example stops comparing software services independently because the assistant has become a reliable source of shortlists. A false generalization in the memory now shapes what the researcher gets a chance to consider. It can narrow deliberation before any visible recommendation is accepted.

There is a further complication. An AI memory system need not behave like a notebook that simply returns an entry. It may summarize, combine, rank, or infer. Barandiaran and Pérez-Verdugo (2024) argue that generative technologies can participate in the formation of intentions in ways that exceed a straightforward model of passive tool use. Their account of generative “midtended” cognition directs attention to the productive contribution of the system and to asymmetries within the relationship.

For the case at hand, the important implication is that the record used tomorrow may already contain an interpretation introduced by the assistant today. If the user later accepts a recommendation based on that interpretation, the acceptance may itself be stored as supporting evidence. A mistaken generalization can acquire an apparently stronger history through the very behavior it helped produce. This is a possible failure mechanism, not a claim that all memory systems exhibit it.

Privacy protections alone would not resolve this problem. The entire cycle could take place on a machine controlled by the user, with no disclosure to another party. What matters is the route from a record to an attributed commitment and from that commitment to action. The person needs some authority over that route, even when they have chosen to delegate most routine operations.

4. What effective revision requires

Three features follow from the argument: access to the grounds of an attribution, a correction that changes subsequent use, and a practical ability to continue elsewhere when the relationship no longer works. They address different ways in which nominal authority can fail.

Start with the grounds. If the assistant says a proposed action conflicts with a standing preference, it should be possible to identify the record supporting that claim. The relevant explanation may be quite simple: a dated decision, its reason, and a note that the broader preference was inferred. This does not require displaying every internal calculation. It requires making the personal premise open to examination. Otherwise, the user is asked to dispute an account whose evidence is inaccessible.

Next comes the effect of correction. A useful record should distinguish an earlier decision from a present commitment. Superseding a decision can preserve the fact that it was made while preventing it from silently governing future action. Deletion can also be appropriate. Neither method is sufficient if obsolete material continues to circulate through summaries, cached context, or other clients. The relevant outcome is that future assistance treats the correction as consequential.

The public memory-hub project illustrates some of these distinctions at the level of data design. Its decision records include reasons and links to superseded decisions. Its import code preserves identifiers and those links when loading an export into another hub (Sami, n.d.). These are useful facilities, but the code does not establish that every connected assistant will interpret the records correctly. A preserved field is an opportunity for better reasoning, not evidence that the reasoning occurred.

Finally, consider departure. If correcting the relationship requires abandoning years of usable context, the user faces a cost created partly by successful dependence on the system. Leaving is formally possible, but it may require rebuilding the very background that made the assistant useful. This supplies a reason for usable transfer: the ability to take enough intelligible context elsewhere to continue important work.

An export containing a long transcript may satisfy a request for data while failing this practical purpose. For the researcher, the decisive information is that a cost restriction belonged to one project and has since been lifted. A receiving assistant needs to distinguish that restriction from a general preference. Transfer therefore concerns meaning and status as well as possession of bytes. Reasons, scope, source, and subsequent corrections can matter more than the sheer volume of material retained.

Perfect reproduction is too demanding. Another model may interpret the same record differently. Some behavioral changes will be improvements. The appropriate aim is a reasonable opportunity to resume the person's projects with their significant commitments and corrections intelligible, rather than to duplicate every previous response. This is a standard for assessing continuity, not a guarantee supplied by any file format.

These features support one another. An inspectable record with no effective correction leaves the user informed but powerless. Correction without a traceable basis can become an unexplained contest between today's instruction and yesterday's summary. Transfer without either can carry the same distortion into a new system. A strong claim to personal assistance must concern the relationship among these capacities.

5. Objections and limits

The first objection is that a good assistant sometimes knows more about a person's patterns than the person presently acknowledges. A researcher may repeatedly overspend and then dismiss every warning as based on outdated information. Would giving the user authority over memory turn the assistant into an agreeable instrument of self-deception?

It need not. Authority over the practical force of a personal attribution is different from authority to make its historical content true or false. The assistant can retain evidence that spending occurred, point out a pattern, and ask whether an exception is intended. The user may also establish a prior rule requiring additional confirmation before changing a budget. What must remain possible is an intelligible route for revising the rule. A commitment can be durable without being irrevocable merely because an assistant remembers it.

This distinction also limits deletion. A person's interest in changing how a memory guides private assistance does not establish an unrestricted entitlement to erase records involving other people's legitimate claims. The argument concerns the assistant's authority to act for the user; it does not supply a general theory of record ownership. Calling data “mine” cannot settle every dispute about its retention.

A second objection is that most people do not want to manage databases. A highly usable hosted service may support their agency better than a self-hosted tool they cannot maintain. This objection is correct against a hardware-based account. Control is effective only when the person can exercise it, directly or through a trusted intermediary. A simple interface for correction may provide more practical authority than administrator access that requires expertise the user lacks. Hosting arrangements should be judged by what they enable, rather than treated as decisive in themselves.

A third objection challenges the emphasis on transfer. A deeply integrated assistant may be valuable precisely because its memory depends on specialized software. Demanding portability could reduce that value, and no replacement may interpret the history equally well. This limits how demanding the requirement can be. The argument does not establish a right to identical performance on any replacement. It establishes a reason to preserve intelligible context wherever feasible and to make losses visible. Users can knowingly accept some dependence; opaque or needlessly imposed dependence is harder to reconcile with continuing authority.

A final objection concerns whether any of this is new. People have long corrected address books, moved documents between computers, and dismissed human assistants. The underlying interests are not unique to AI. The distinctive pressure comes from combining a changing personal record with a system that interprets it and repeatedly uses those interpretations to act. The contribution here is to connect correction and transfer through that delegation relationship. A memory system deserves scrutiny not only for what it stores, but for how stored material becomes a reason to act for someone.

6. Continuity without freezing the person

There remains a temptation to identify a personal AI with a persistent character: the same name, voice, conversational manner, or apparent sense of familiarity. Those features can matter to a user. They are poor substitutes for continuity of the user's projects and authority. An assistant could keep its recognizable voice after losing the reasons behind an important decision. Another could sound quite different while preserving those reasons and responding correctly to a later revision.

Nothing in this account establishes that a model with access to a continuing record is the same conscious subject across sessions. That is a separate philosophical question. The continuity at issue here is practical: important work can proceed, relevant history remains available, and the person can change what that history is allowed to decide. Conflating this with persistence of a conscious self would make the argument depend on a claim it does not need.

The same caution applies to inheritance, a possibility raised in “Two AIs.” A person may leave a valuable record of decisions and reasons to someone else. The usefulness of that record would not demonstrate the survival of the original person. Nor would inheriting it give the successor an obligation to continue the same preferences. If personal memory is valuable because it supports agency, a successor must have room to interpret, refuse, and revise what has been received.

The attraction of personal AI begins with a reasonable wish: not to explain one's life and work from the beginning every time a tool changes. But being remembered has two sides. A remembered reason can keep a project coherent; a remembered interpretation can keep returning a person to a position they have left. More memory increases what can be carried forward. It does not decide what deserves to govern the next step.

A defensible account of personal assistance must therefore make room for change. The assistant may preserve the record and question an inconsistency. It should also let the person say what the inconsistency means. For an AI to count as meaningfully personal in the sense defended here, the history it uses must remain open to the continuing authority of the person whose work and life it is meant to assist.

References

- Barandiaran, Xabier E., and Marta Pérez-Verdugo. 2024. “Generative Midtended Cognition and Artificial Intelligence. Thining with Thining Things.” arXiv:2411.06812, version 1.
<https://doi.org/10.48550/arXiv.2411.06812>
- Clark, Andy, and David J. Chalmers. 1998. “The Extended Mind.” *Analysis* 58 (1): 7-19.
<https://doi.org/10.1093/analys/58.1.7>
- Gabriel, Iason, et al. 2024. “The Ethics of Advanced AI Assistants.” arXiv:2404.16244, version 2.
<https://doi.org/10.48550/arXiv.2404.16244>
- Sami, Zawwar. 2026. “Two AIs.” Zawwar's Substack, May 22. <https://zawwar16.substack.com/p/two-ais>
- Sami, Zawwar. n.d. zai-memory-hub. Software repository, README and scripts/import_export.py. GitHub. Accessed October 2, 2026. <https://github.com/Zawwarsami16/zai-memory-hub>